

---


# CIS 721 - Real-Time Systems

## Lecture 9: Preemption Thresholds

---

Mitch Neilsen  
**neilsen@ksu.edu**

# Daily Humor



# Outline

- Clock-Driven Scheduling (Ch. 5)
- **Priority-Driven Scheduling**
  - **Periodic Tasks (Ch. 6)**
    - Optimal Priority Assignment
    - Arbitrary Start Times
      - Leung's Feasibility Test
      - Audsley's Feasibility Test
    - **Arbitrary Deadlines**
    - **Preemption Thresholds**
  - Aperiodic and Sporadic Tasks (Ch. 7)

# Terms and Concepts

- A **task set**  $\Gamma = \{\tau_1, \tau_2, \dots, \tau_n\}$  is a collection of related tasks.
- Each **periodic task**  $\tau_i$  is characterized by:
  - an **execution time** or **run-time** (  $C_i$  ),
  - a **period** (  $T_i$  ),
  - a **(relative) deadline** (  $D_i$  ), and
  - a **phase** or **offset** (  $\varphi_i$  or  $O_i$  ).

# Response Time Analysis

- For each (potentially overlapping) release, a worst-case completion time  $w_i(q)$  is defined by:

$$w_i^{n+1}(q) = q \cdot C_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n(q)}{T_j} \right\rceil \cdot C_j$$

$$w_i^0(q) = C_i + (q-1) \cdot T_i$$

where  $q$  is the instance or job number and  $w_i(q)$  is the least fixed point of  $w_i^n(q)$ .

- The response time of the  $q^{th}$  instance,  $R_i(q)$ , is given by  $R_i(q) = w_i(q) - (q-1) T_i$ .

# Response Time Analysis (cont.)


- Set  $q' = \min \{ q \mid R_i(q) \leq T_i \}$ .
- Then, the **level-i busy period** is  $[ 0, w_i(q') ]$ .
- The **worst-case response time**  $R_i$  is given by:

$$R_i = \max_{q=1,2,\dots,q'} \{R_i(q)\}$$

- If  $R_i \leq D_i$  for all  $i$ , the system is **schedulable**.

# Lehoczky's Example

<b>Task</b>	<b>Period</b>	<b>Run-Time</b>	<b>Phase</b>	<b>Deadline</b>
$\tau_i$	$T_i$	$C_i$	$\phi_i$	$D_i$
<hr/>				
$\tau_1$	<b>70</b>	<b>26</b>	<b>0</b>	<b>68</b>
$\tau_2$	<b>100</b>	<b>62</b>	<b>0</b>	<b>118</b>





# General Time-Driven Analysis

- Check to see if the first job completes before it's deadline and before the second job in the same task is released.
- If it completes before it's deadline, but not before the second job is released, then check the second job.
- In general, continue to check all jobs over a level-i busy period; that is, until a deadline is missed or until one job completes before the next one is released.

# Preemption Thresholds

- Y. Wang and M. Saksena, “Scheduling Fixed-Priority Tasks with Preemption Threshold”, In Proceedings of the IEEE Intl. Conf. on Real-Time Computing Systems and Applications, Dec. 1999.
- Scheduling with Preemption Thresholds
  - Task Model and Run-Time Model
  - Response Time Analysis
  - Priority and Preemption Threshold Assignment Algorithms
  - Example: ThreadX Real-Time Operating System

# Task Model

- Task Set  $\Gamma = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$ 
  - Each task  $\tau_i$  is characterized by  $(C_i, T_i, D_i)$ , denoted  $\tau_i \sim (C_i, T_i, D_i)$ .
  - Each task  $\tau_i$  is assigned a priority  $\pi_i \in \{1, 2, \dots, n\}$
  - and a preemption threshold  $\gamma_i \in \{\pi_i, \pi_i + 1, \dots, n\}$ .
- Notes:
  - 1 = lowest priority, n = highest priority.
  - $\pi_i$  = static priority.
  - $\gamma_i$  = dynamic priority.

# Run-Time Model

- Modified fixed-priority, preemptive scheduling.
- When task  $\tau_i$  is released, it is scheduled using its static priority  $\pi_i$ .
- After task  $\tau_i$  starts executing, another task  $\tau_j$  can preempt  $\tau_i$  only if  $\pi_j > \gamma_i \geq \pi_i$ .


# Extremes

- If  $\gamma_i = \pi_i$  for each  $i$ , then the result is **preemptive**, priority-based scheduling.
- If  $\gamma_i = n$  (max. priority) for each  $i$ , then the result is **non-preemptive**, priority-based scheduling.

# Example

Task	$C_i$	$T_i$	$D_i$	$\pi_i$	WCRT Preemptive	WCRT Non-Preemptive
$\tau_1$	20	70	50	3	20	<b>55</b>
$\tau_2$	20	80	80	2	40	75
$\tau_3$	35	200	100	1	<b>115</b>	75

Task	Priority	Preemption Threshold	WCRT
$\tau_1$	3	3	40
$\tau_2$	2	3	75
$\tau_3$	1	2	95



**Figure 1. Run-time Behavior with Preemption Threshold**

# Problem Statement

- Given a task set  $\Gamma = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$ , determine if there exists an assignment  $\{(\pi_i, \gamma_i) \mid i = 1, 2, \dots, n\}$  such that  $\Gamma$  is schedulable.
- In other words, determine if there exists an **optimal** assignment of task priorities and preemption thresholds.



# Solutions

- **Brute Force** - try all possible assignments of priorities and preemption thresholds.
  - Time Complexity in  $O(n! n!)$   $\Rightarrow$  not feasible for large  $n$ .
- Use a **Branch and Bound Algorithm** to perform an efficient search for priorities and preemption thresholds.

# Three Step Process


## ■ Response Time Analysis

- Given assignment  $\{ (\pi_i, \gamma_i) \mid i = 1, 2, \dots, n \}$ , compute the worst-case response time ( $R_i$  or  $WCRT_i$ ) for each task  $\tau_i$ .
- A task set  $\Gamma$  is schedulable iff  $R_i \leq D_i$  for all  $i$ .
- Given a priority assignment  $\{ \pi_i \mid i = 1, \dots, n \}$ , determine a feasible set of **preemption thresholds**, if such a set exists.
- Use a branch and bound algorithm to search for a feasible assignment set of **priorities** (and preemption thresholds).

# Response Time Analysis


- The **blocking time** of task  $\tau_i$  is denoted  $B(\tau_i)$ . Blocking occurs if a lower priority task is running and task  $\tau_i$  cannot preempt it.

$$B(\tau_i) = \max_j \{C_j / \gamma_j \mid \pi_i > \pi_j\}$$



# Busy Period Analysis

- A **critical instant** occurs when all higher priority tasks arrive at the same time, and the task that contributes to the maximum blocking arrives at the critical instant -  $\varepsilon$ .



# Divide Busy Period

- Divide the busy period for  $\tau_i$  into two parts:
  - the length of time from the critical instant (time 0) to the point when  $\tau_i$  starts executing its  $q^{\text{th}}$  job (  $\mathbf{S}_i(\mathbf{q})$  ).
  - the length of time from the time  $\tau_i$  starts executing its  $q^{\text{th}}$  job until it finishes executing its  $q^{\text{th}}$  job (  $\mathbf{F}_i(\mathbf{q}) - \mathbf{S}_i(\mathbf{q})$  ).
- Let  $q = 1, 2, \dots, m$  until we reach  $q = m$  s.t.  $F_i(m) \leq m T_i$  that is, the  $m^{\text{th}}$  job completes before the next job is released.
- Then,

$$R_i = \max_{q \in \{1, \dots, m\}} \{ F_i(q) - (q - 1)T_i \}$$

# Worst-Case Start Time ( $S_i(q)$ )

$$S_i(q) = B(\tau_i) + (q-1)C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \pi_i}} (1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor) C_j$$

# Worst-Case Finish Time ( $F_i(q)$ )

$$F_i(q) = S_i(q) + C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \gamma_i}} \left( \left\lceil \frac{F_i(q)}{T_j} \right\rceil - \left( 1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor \right) \right) C_j$$

# $WCRT(\pi_i, \gamma_i)$

Algorithm to compute  $R_i$

Input :  $C_1, \dots, C_m, T_1, \dots, T_m, \pi_1, \dots, \pi_m, \gamma_1, \dots, \gamma_m$

Output :  $R_1, R_2, \dots, R_m$

done = FALSE

q = 1

while (not done)

    compute  $S_i(q)$  and  $F_i(q)$

    if  $F_i(q) \leq q T_i$  then

        done = TRUE

        m = q

    else

        q = q + 1

    end if

end while

$R_i = \max_{q \in \{1, \dots, m\}} (F_i(q) - (q - 1) T_i)$



# Example

Task	$C_i$	$T_i$	$D_i$	$\pi_i$	WCRT Preemptive	WCRT Non-Preemptive
$\tau_1$	20	70	50	3	20	<b>55</b>
$\tau_2$	20	80	80	2	40	75
$\tau_3$	35	200	100	1	<b>115</b>	75

Task	Priority	Preemption Threshold	WCRT
$\tau_1$	3	3	40
$\tau_2$	2	3	75
$\tau_3$	1	2	95

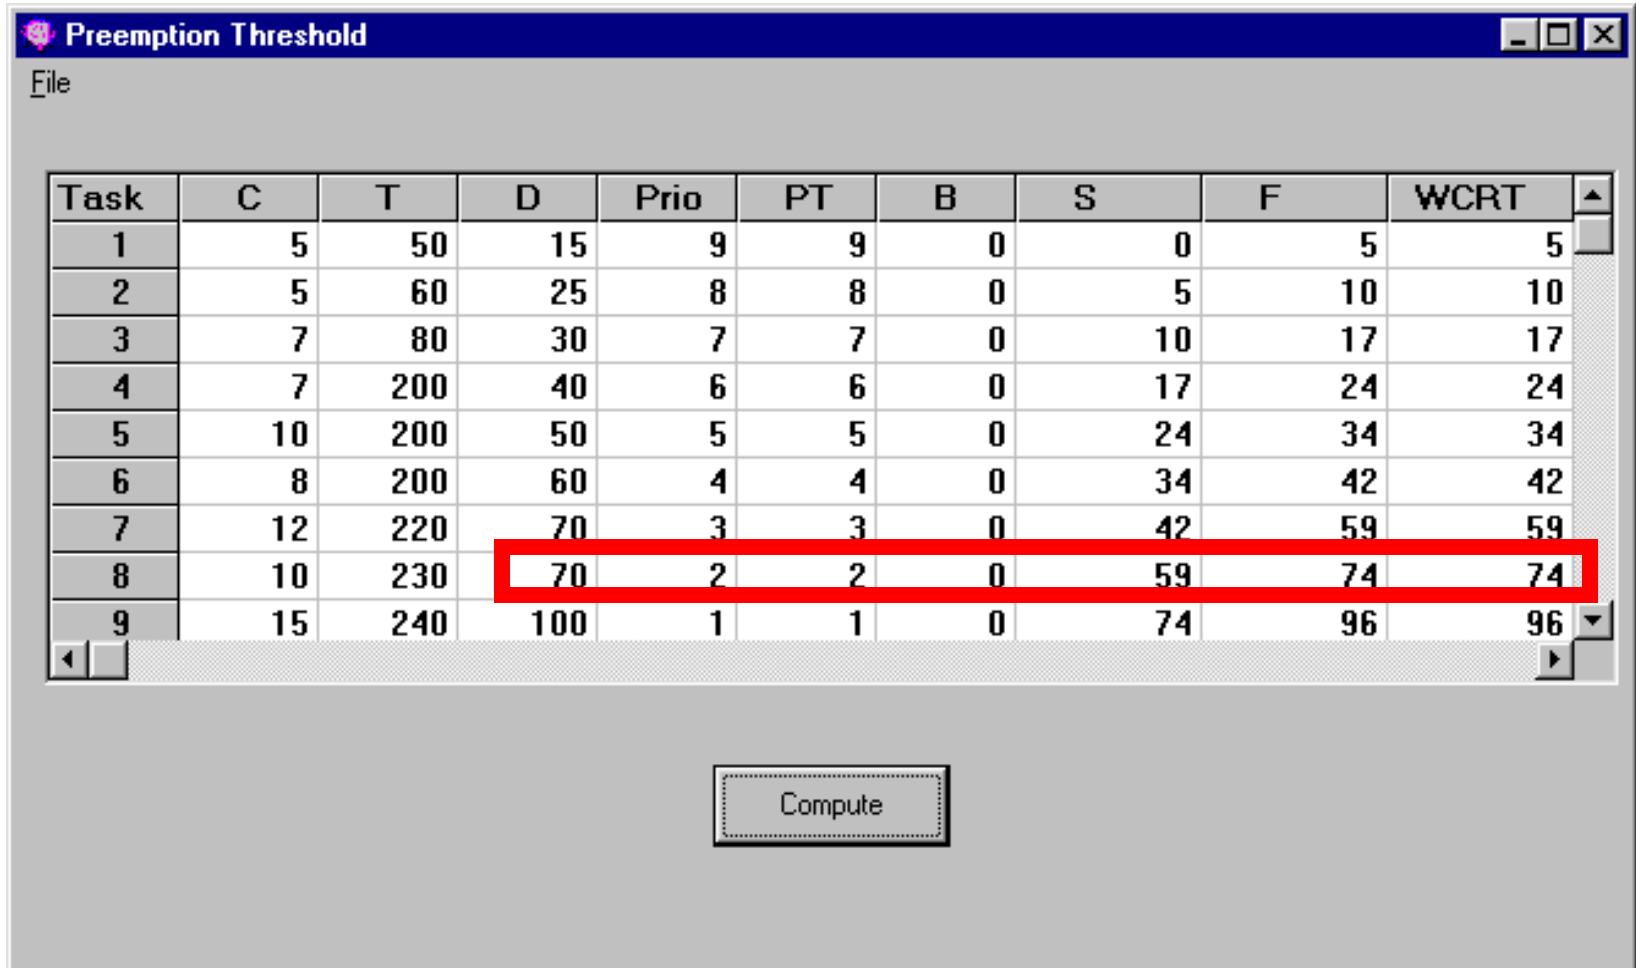
# Preemption Threshold Assignment

- **Lemma 5.1:** Changing the preemption threshold of task  $\tau_i$  from  $\gamma_1$  to  $\gamma_2$  may only affect the worst-case response time of task  $\tau_i$  and those tasks whose priority is between  $\gamma_1$  and  $\gamma_2$ .
- **Corollary 5.1:** The worst-case response time of task  $\tau_i$  will not be affected by the preemption threshold assignment of any higher priority task; e.g., any task  $\tau_j$  with  $\pi_j > \pi_i$ .
- **This implies that we should assign preemption thresholds from lowest to highest priority.**

# Preemption Threshold

- **Theorem 5.1:** Start with a schedulable system with  $n$  tasks. If decreasing the value of  $\gamma_j$  does not change the schedulability of task  $\tau_j$ , then the whole system is still schedulable.
- **Idea:** Keep  $\gamma_j$  as small as possible for each task.
- **Lemma 5.2:** (Quick Test) If setting  $\gamma_j = n$  cannot make task  $\tau_j$  schedulable, then the task set is not schedulable.

# Preemptive Scheduling

A screenshot of a software window titled "Preemption Threshold". The window has a menu bar with "File" and a table of task scheduling data. The table has columns: Task, C, T, D, Prio, PT, B, S, F, and WCRT. Row 8 is highlighted with a red rectangle. Below the table is a "Compute" button.

Task	C	T	D	Prio	PT	B	S	F	WCRT
1	5	50	15	9	9	0	0	5	5
2	5	60	25	8	8	0	5	10	10
3	7	80	30	7	7	0	10	17	17
4	7	200	40	6	6	0	17	24	24
5	10	200	50	5	5	0	24	34	34
6	8	200	60	4	4	0	34	42	42
7	12	220	70	3	3	0	42	59	59
8	10	230	70	2	2	0	59	74	74
9	15	240	100	1	1	0	74	96	96

Compute

# Non-Preemptive Scheduling

Preemption Threshold

File

Task	C	T	D	Prio	PT	B	S	F	WCRT
1	5	50	15	9	9	15	15	20	20
2	5	60	25	8	9	15	20	25	25
3	7	80	30	7	9	15	25	32	32
4	7	200	40	6	9	15	32	39	39
5	10	200	50	5	9	15	39	49	49
6	8	200	60	4	9	15	49	57	57
7	12	220	70	3	9	15	67	79	79
8	10	230	70	2	9	15	79	89	89
9	15	240	100	1	9	0	74	89	89

Compute

## Algorithm: Assign Preemption Thresholds

*// Assumes that task priorities are already known*

- (1) **for** ( $i := 1$  to  $n$ )
- (2)      $\gamma_i = \pi_i$   
      *// Calculate worst-case response time of  $\tau_i$*
- (3)      $\mathcal{R}_i = \text{WCRT}(\tau_i, \gamma_i)$  ;
- (4)     **while** ( $\mathcal{R}_i > D_i$ ) **do**     *// while not schedulable*
- (5)          $\gamma_i++$  ; *// increase threshold*
- (6)         **if**  $\gamma_i > n$  **then**
- (7)             **return** FAIL; *// system not schedulable.*
- (8)         **endif**
- (9)          $\mathcal{R}_i = \text{WCRT}(\tau_i, \gamma_i)$  ;
- (10)    **end**
- (11) **end**
- (12) **return** SUCCESS



## Preemption Threshold



File

Task	C	T	D	Prio	PT	B	S	F	WCRT	▲
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	0	10	17	17	
4	7	200	40	6	6	0	17	24	24	
5	10	200	50	5	5	0	24	34	34	
6	8	200	60	4	4	0	34	42	42	
7	12	220	70	3	3	0	42	59	59	
8	10	230	70	2	2	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	▼

Compute

Task	C	T	D	Prio	PT	B	S	F	WCRT	
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	0	10	17	17	
4	7	200	40	6	6	0	17	24	24	
5	10	200	50	5	5	0	24	34	34	
6	8	200	60	4	4	0	34	42	42	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	3	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	

Compute



Task	C	T	D	Prio	PT	B	S	F	WCRT	
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	0	10	17	17	
4	7	200	40	6	6	0	17	24	24	
5	10	200	50	5	5	0	24	34	34	
6	8	200	60	4	4	10	44	57	57	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	4	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	

Compute

Task	C	T	D	Prio	PT	B	S	F	WCRT	
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	0	10	17	17	
4	7	200	40	6	6	0	17	24	24	
5	10	200	50	5	5	10	34	44	44	
6	8	200	60	4	4	10	44	57	57	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	5	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	

Compute

Task	C	T	D	Prio	PT	B	S	F	WCRT	▲
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	0	10	17	17	
4	7	200	40	6	6	10	27	34	34	
5	10	200	50	5	5	10	34	44	44	
6	8	200	60	4	4	10	44	57	57	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	6	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	▼

Compute



# Preemption Threshold



File

Task	C	T	D	Prio	PT	B	S	F	WCRT	
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	0	5	10	10	
3	7	80	30	7	7	10	20	27	27	
4	7	200	40	6	6	10	27	34	34	
5	10	200	50	5	5	10	34	44	44	
6	8	200	60	4	4	10	44	57	57	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	7	0	59	74	74	
9	15	240	100	1	1	0	74	96	96	

Compute

Task	C	T	D	Prio	PT	B	S	F	WCRT	▲
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	10	15	20	20	
3	7	80	30	7	7	10	20	27	27	
4	7	200	40	6	6	10	27	34	34	
5	10	200	50	5	5	10	34	44	44	
6	8	200	60	4	4	10	44	57	57	
7	12	220	70	3	3	10	57	74	74	
8	10	230	70	2	8	0	59	69	69	
9	15	240	100	1	1	0	74	96	96	▼

◀
▶

Compute

File

Task	C	T	D	Prio	PT	B	S	F	WCRT	▲
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	10	15	20	20	
3	7	80	30	7	7	12	22	29	29	
4	7	200	40	6	6	12	29	36	36	
5	10	200	50	5	5	12	36	46	46	
6	8	200	60	4	4	12	46	59	59	
7	12	220	70	3	7	10	57	74	74	
8	10	230	70	2	8	0	59	69	69	
9	15	240	100	1	1	0	74	96	96	▼

Compute

File

Task	C	T	D	Prio	PT	B	S	F	WCRT	▲
1	5	50	15	9	9	0	0	5	5	
2	5	60	25	8	8	12	17	22	22	
3	7	80	30	7	7	12	22	29	29	
4	7	200	40	6	6	12	29	36	36	
5	10	200	50	5	5	12	36	46	46	
6	8	200	60	4	4	12	46	59	59	
7	12	220	70	3	8	10	57	69	69	
8	10	230	70	2	8	0	59	69	69	
9	15	240	100	1	1	0	74	96	96	▼

Compute

# Preemption Thresholds

The screenshot shows a window titled "Preemption Threshold". Below the title bar is a menu bar with "File". The main area contains a table with 10 rows and 10 columns. The columns are labeled Task, C, T, D, Prio, PT, B, S, F, and WCRT. The rows contain numerical values. Two vertical green rectangles highlight the 'D' column and the 'WCRT' column. At the bottom center is a button labeled "Compute".

Task	C	T	D	Prio	PT	B	S	F	WCRT
2	5	60	25	8	9	12	17	22	22
3	7	80	30	7	8	12	22	29	29
4	7	200	40	6	7	12	29	36	36
5	10	200	50	5	6	12	36	46	46
6	8	200	60	4	5	12	46	59	59
7	12	220	70	3	8	10	57	69	69
8	10	230	70	2	8	0	59	69	69
9	15	240	100	1	1	0	74	96	96
10									

Compute



# Preemption Thresholds – ThreadX RTOS

- Response Time Analysis to compute WCRT given Priorities and Preemption Thresholds
- Algorithm to optimally assign Preemption Thresholds given Priority Assignment
- **Algorithm to Assign Priorities**


---

# Finding Priority Assignment

- **Problem:** How to find an optimal assignment of task priorities and preemption thresholds.
  - **Solution Proposed in Paper:**
    - Arrange tasks into two sets -- unsorted (remaining higher priority tasks) and sorted (lower priority tasks).
    - Recursively add tasks from unsorted list to sorted list based on “**lateness**” heuristic.
    - Tasks are added in priority order, from lowest to highest priority.
-

# Priorities and Preemption Thresholds

- Search  $(\{\tau_1, \dots, \tau_n\}, 1)$



- Initially determine if the task set is schedulable using preemptive priority-based scheduling without preemption thresholds (e.g., can priorities be assigned using RM?).

---

# “Lateness” Heuristic

- From the unsorted list, select the task with the smallest lateness, and add it to the sorted list.
  - Since the task selected has the smallest “lateness” (delay), it should need a lower priority and smaller preemption threshold, leaving more time for higher priority tasks.
-

# Greedy Assignment Algorithm

**Algorithm: GreedyAssignment(RemainingTasks, nextPriority)**

*/\* Terminating Condition \*/*

- (1)   **if** (RemainingTasks == NULL) **then**
  - /\* Call the algorithm in Figure 1 for optimal preemption threshold assignment \*/*
- (2)       **if** (AssignThresholds() == SUCCESS) **then return** SUCCESS
- (3)       **else return** FAIL
- (4)       **endif**
- (5)   **endif**

# Greedy Assignment Algorithm (cont.)

*/\* Assign Heuristic Value to Each Task \*/*

- (6)   **foreach**  $\tau_k$  in RemainingTasks **do**
- (7)        $\pi_k := \text{nextPriority}$  ;   */\* tentative assignment \*/*
- (8)        $\mathcal{R}_k := \text{WCRT}(\tau_k)$ ;   */\* compute response time \*/*
- (9)       **if**  $\mathcal{R}_k \geq D_k$  **then**  $h\_val_k := D_k - \mathcal{R}_k$
- (10)      **else**  $h\_val_k := \text{GetBlockingLimit}(\tau_k)$  ;
- (11)      **endif**
- (12)       $\pi_k := \pi$  ;   */\* reset, to allow computing heuristic value for other tasks \*/*
- (13)   **end**

# Get Blocking Limit Function

Input :  $\tau_k, D_k$

Output : Blocking limit of  $\tau_k$

$R_k = WCRT(\tau_k)$

$Max = D_k - R_k$

$Limit = 0$

**For**  $B(\tau_k) = 1$  **to**  $Max$

$R_k = WCRT(\tau_k, B(\tau_k))$

**If**  $R_k > D_k$  **Then** *Break*

**Else**  $Limit = B(\tau_k)$

**End For**

**Return**  $Limit$

# Greedy Assignment Algorithm (cont.)

```
(14)  /* Select the task with the largest heuristic value next */
(15)   $\tau_k := \text{max\_heuristic\_val}(\text{RemainingTasks})$  ;
(16)   $\pi_k := \text{nextPriority}$  ;    /* final priority assignment */
(17)  /* Recursively Assign Priorities to Remaining Tasks */
(18)  if GreedyAssignment(RemainingTasks  $\leftarrow \tau_k, \text{nextPriority}+1$ ) == SUCCESS then
(19)      return SUCCESS ;
(20)  return FAIL ;
```

---



# Note

- There are cases when this heuristic algorithm is not able to find a feasible assignment, even though a non-preemptive priority assignment algorithm is able to find a solution.
- Thus, we could try a non-preemptive assignment algorithm first, before using this heuristic algorithm (or use a better algorithm)

---

# Depth-First Search

- Perform a depth-first search to find an optimal priority assignment.
  - When a leaf is reached, call `AssignThresholds( )` to see if an optimal preemption threshold assignment exists; if not, continue searching.
-

---

# Summary

- Preemption thresholds provide a way of generalizing both preemptive and non-preemptive scheduling in a single framework.
  - Read Y. Wang and M. Saksena's paper on preemption thresholds.
-